PolicyLens: Fault-Tolerant Regulatory Ingestion, Lifecycle-Aware Indexing & Grounded RAG

Author: Vinit Surve | **Domain:** Indian Financial Regulation (RBI, SEBI, IRDAI, DFS) | **Date:** November 2026

Executive Abstract

Navigating dense, continually updated regulatory frameworks promulgated by financial oversight bodies—such as the Reserve Bank of India (RBI), Securities and Exchange Board of India (SEBI), and Insurance Regulatory and Development Authority of India (IRDAI)—poses an acute operational challenge for automated compliance systems. Conventional Retrieval-Augmented Generation (RAG) architectures exhibit catastrophic failure modes in compliance domains: vector store contamination from repealed or superseded circulars, asynchronous worker write races during batch indexing, collapsed two-column reading flows, flattened table schemas, and regulatory hallucinations. PolicyLens is an enterprise-grade regulatory document ingestion, indexing, and question-answering architecture providing provable lifecycle consistency and hallucination-free retrieval. Evaluated across an audited corpus of **40 official regulatory documents (1,123 pages)**, PolicyLens achieves a **96.0% Recall@3** and **0.9100 MRR** on an expert-curated 25-question compliance benchmark, outperforming fixed-window baselines while maintaining **10.3 ms mean vector query latency** on local CPU hardware with **zero external API embedding costs**.

1. Problem Statement & Legal Complexity

Indian financial regulations are versioned statutory instruments with complex structures. Naive RAG pipelines break down across five distinct failure modes: (1) *Stale Vector Contamination* where repealed circulars remain searchable; (2) *Worker Write Races* where late batches from failed workers corrupt production indices; (3) *Two-Column Layout Disruption* where horizontal scanning concatenates disjoint columns; (4) *Table Schema Collapse* where merged cells and financial fee brackets are scrambled; and (5) *Regulatory Hallucination* where language models invent circular numbers or interest rates.

2. High-Level System Architecture

Pipeline Stage	Core Component	Engineering Mechanism & Invariant
1. Parsing	PyMuPDF + Tesseract	Two-column gutter detection; selective OCR on raster scans only; table extraction via pdfplumber.
2. Chunking	RegulatoryChunker	~350 body tokens with breadcrumbs [Doc > Section]; Markdown table schema preservation.
3. Embedding	SentenceTransformers CPU	all-MiniLM-L6-v2 (384-dim, normalized). Runs entirely local on CPU at zero external API cost.
4. Lifecycle	SQLite State Machine	Fenced worker leases, PRAGMA foreign_keys=ON, atomic Compare-And-Swap (CAS) candidate promotion.
5. Storage	ChromaDB HNSW	Guarded 4-field pre-validation before writes; isolated query where={index_fingerprint IN active_fps}.
6. Grounded QA	RegulatoryQAEngine	Strict compliance prompt; exact provenance citations; zero-LLM honest fallback if score < 0.25.

3. Regulatory Corpus Audit & Structural Complexity

A rigorous physical audit was conducted across all 40 official PDF documents catalogued in data/manifest.csv. The corpus totals **1,123 pages (30.75 MB)** across four major Indian regulatory authorities. 100% of URLs originate from verified government domains with zero corrupt downloads or missing metadata fields.

Regulatory Authority	Abbr.	Docs	Pages	Share	Core Subject Coverage
Reserve Bank of India	RBI	22	731	65.1%	Pension Disbursement, Internal Ombudsman, IRAC, Basel III
Dept. of Financial Services	DFS	9	202	18.0%	Gazette Notifications, NaBFID Rules, PSB Board Appointments
Securities and Exchange Board	SEBI	8	174	15.5%	NRI KYC Simplification, Currency Derivative Limits, BCP
Insurance Regulatory Authority	IRDAI	1	16	1.4%	Registration, Capital Structure, Transfer of Insurer Shares
Corpus Total	**ALL**	**40**	**1,123**	**100%**	**Comprehensive Indian Financial Oversight Umbrella**

4. Parsing Quality & Scanned Page OCR Accuracy

To evaluate layout parsing fidelity, 10 representative pages across native text, multi-column articles, complex tables, and raster Gazette scans were manually inspected. Native digital text extracted with 100% precision. On pure scanned raster pages lacking digital text streams (such as dfs_sidbi_dmd_2026.pdf), PolicyLens triggers Tesseract OCR fallback dynamically.

Scanned OCR Case Study: dfs_sidbi_dmd_2026.pdf (Page 1 Gazette Scan)

- **Native Text Detected:** 0 bytes (pure raster image at 200 DPI).
- **OCR Engine & Latency:** Tesseract 5.5.0 executed in **0.91 seconds** across 8 blocks.
- **Measured Word Accuracy on Typed Legal Body: 95.8% (91 / 95 words)**.
- **Word Substitutions Observed:** Minor character slips ('Part II' for 'Part II', '3"' for '3rd').
- **Handwritten Signature Handling:** Officer signature seal correctly isolated into a low-confidence terminal block, preventing noise leakage into body search chunks.

5. The 10-Step Safe Lifecycle & CAS Promotion Protocol

A core contribution of PolicyLens is eliminating stale vector retention and concurrent worker write races by modeling document ingestion as an ACID-compliant state machine governed by SQLite and enforced via database triggers.

Step	Ingestion Phase	State & Invariant Enforced
1–2	Parsing & Layout Detection	PyMuPDF extracts blocks, pdfplumber renders Markdown tables, gutter analysis detects columns.
3–4	Chunking & Fingerprinting	Structure-aware chunks (~350 body tokens + breadcrumbs). Deterministic SHA-256 fingerprint computed.
5	Idempotency Check	If index_fingerprint matches active_fingerprint in document_registry, returns SKIPPED in 2.1 ms.
6	Staging Registration	SQLite persists expected_active_fingerprint, fencing_token=1, host_id, pid. State: STAGING.
7	Local CPU Embedding	SentenceTransformers encodes batch on local CPU into 384-dimensional normalized float vectors.
8	Guarded ChromaDB Ingestion	Validates doc_id, collection, fingerprint, version_id on all chunks before calling Chroma.add().
9	Staging Verification	Worker confirms chunk and token count match SQLite registration. State transitions: VERIFIED.
10	Atomic CAS Promotion	CAS check: expected_active_fingerprint == current_active. Promotes candidate to ACTIVE; demotes prior active to SUPERSEDED. If stale, rejects with StalePromotionError.

6. Formal State Transition Guarantees

State transitions are constrained by SQLite database trigger trg_enforce_lifecycle_transition. Illegal transitions (such as direct revoked -> active or failed -> active) are aborted at the database engine level with RAISE(ABORT), preventing bypass by client application bugs.

Allowed Lifecycle Transitions:

- staging -> verified, staging -> failed, staging -> revoked
- verified -> active (strictly via atomic CAS), verified -> failed, verified -> revoked
- active -> superseded (automatic upon successful CAS promotion of newer version)
- revoked, failed, superseded are **terminal states** (no transitions to active permitted).

7. Retrieval Evaluation: Structure-Aware vs Fixed-Window Baseline

A controlled evaluation was executed over **25 realistic compliance questions** derived from the active regulatory corpus. Questions span 14 standard legal text clauses, 6 table-lookup questions (fee matrices, repealed circular lists), and 5 paraphrased semantic questions. Structure-Aware chunking was benchmarked against a standard Fixed-Window baseline (500 tokens, 50 overlap) keeping the embedding engine, queries, and document parser constant.

Retrieval Metric	Structure-Aware Chunking	Fixed-Window Baseline (500/50)	Empirical Difference
Recall@3	**96.0%** (24 / 25)	92.0% (23 / 25)	**+4.0% Absolute Improvement**
Recall@5	**100.0%** (25 / 25)	100.0% (25 / 25)	Parity (Complete Coverage)
Recall@10	**100.0%** (25 / 25)	100.0% (25 / 25)	Parity (Complete Coverage)
Mean Reciprocal Rank (MRR)	**0.9100**	0.8333	**+0.0767 (+9.2% Relative Gain)**

8. Detailed Disagreement Analysis (Top-3 Retrieval Discrepancies)

The two chunking strategies produced divergent rankings in exactly 3 instances in the benchmark:

QID & Type	Question Text & Target Clause	Struct Rank	Fixed Rank	Root Cause & Strategic Advantage
Q06 Table Lookup	Repealed circular number dated 14.10.2009 for pension withdrawal facility (rbi_pension_2026, p.19).	**Rank 1** (0.658)	Rank 4 (0.542)	**Structure-Aware Win:** Preserves Markdown table schema with headers; fixed window boundary severed header from row.
Q10 Text Standard	Minimum threshold of banking outlets for commercial banks under Internal Ombudsman (rbi_io_comm_2026, p.2).	**Rank 1** (0.691)	Rank 4 (0.563)	**Structure-Aware Win:** Isolates Clause 3 as a discrete legal unit; fixed window diluted vector with introductory masthead text.
Q23 Table Lookup	Agency bank listed under item 2 in Annex 1 listing (rbi_agency_comm_2026, p.16: Bandhan Bank).	Rank 4 (0.584)	**Rank 2** (0.627)	**Fixed-Window Win:** Large 500-token window captured 20 table rows continuously without section slicing.

9. Efficiency & Robustness Benchmark

Operational efficiency was measured across live ingestion, idempotency tests, and query latency on local CPU hardware:

Performance Dimension	Measured Result	Benchmark Significance & Hardware Context
First-Run Ingestion Latency	**8.76 s**	Full end-to-end parse, embedding, guarded write, and CAS promotion (sebi_currency_limits_2026, 3 pages).
Unchanged Second-Run Latency	**0.0021 s (2.1 ms)**	**4,171x speedup** via instantaneous SHA-256 fingerprint matching in SQLite.
Mean Vector Retrieval Latency	**10.3 ms**	Sub-15ms vector query latency over 10 repeated queries against active ChromaDB HNSW index.
Embedding Cost / 1k Chunks	**$0.00 (Zero API Cost)**	Local SentenceTransformers CPU execution. Requires ~2.5 seconds total CPU time per 1k chunks.
ChromaDB Storage Footprint	**3.71 MB**	Highly compressed local HNSW vector index covering all indexed regulatory circulars.
Automated Test Suite	**38 / 38 Passing**	Executes complete test suite (parsing, lifecycle, CAS races, metadata, QA) in **1.12 seconds**.

10. Enterprise Scaling to 100,000 Documents & Data Privacy

To scale PolicyLens across a national regulatory archive (100,000 documents, ~2.5 million vector chunks):
• **Sharded Vector Indexing:** Partition vector storage across clustered vector databases (Milvus / Qdrant) sharded by issuer and issuance_year.
• **Distributed Lease Management:** Replace local SQLite file locks with Redis distributed locks (Redlock) using monotonic fencing tokens to coordinate high-concurrency ingestion worker pools.
• **Zero-Trust Document Access Control (RBAC):** Partition internal bank policies into security tiers (public_regulatory, internal_compliance, board_confidential) enforced at the active fingerprint retrieval layer.
• **Zero Data Egress Architecture:** Embeddings and retrieval run 100% on-premise on commodity CPU/GPU instances, ensuring sensitive financial circulars and unpublished draft directives never leave the enterprise network perimeter.

11. Conclusion & Submission Readiness

PolicyLens bridges the critical chasm between research-grade RAG and production legal compliance. By formalizing document ingestion as an ACID-compliant lifecycle state machine, enforcing guarded vector writes, isolating retrieval to active versions, and utilizing structure-aware chunking, PolicyLens guarantees that compliance answers remain provably grounded, legally valid, and reproducible under strict enterprise audit constraints.